

VAL CoPilot

Azure AI Foundry Deployment Guide

Part A = Azure Foundry portal steps | Part B = GitHub repo / deploy script

Complete **Part A** fully before running `deploy_to_foundry.py`.

Step-by-step procedure to provision the LangChain cognitive routing agent into Microsoft Azure AI Foundry using `copilot_agent/deploy_to_foundry.py`.

Audience: Azure Cloud Architects / DevSecOps Engineers

SDKs: `azure-ai-projects==1.0.0` · `azure-ai-agents` · `azure-identity`

This guide is split into two parts:

Part	Where you work	Outcome
A — Azure Foundry side	Azure Portal + <code>ai.azure.com</code>	Project, model, RBAC, endpoint ready
B — Code / script side	Developer machine with the GitHub repo	Managed agent created with <code>SYSTEM_PROMPT</code> + tools

Do Part A completely before Part B. The deploy script cannot create the Foundry resource, project, or model deployment for you.

1. Overview

This guide deploys VAL CoPilot as a **managed Azure AI Foundry Agent**. The deployment script authenticates with Entra ID, wraps Fabric SQL and Azure AI Search MCP tools in a Foundry **FunctionTool / ToolSet**, and creates the agent using the existing multi-agent **SYSTEM_PROMPT** (compliance audits, red-flagging, financial exposure projections, and renewal strategy sheets).

What gets provisioned (by the script)

- Managed agent named `val-copilot` (configurable)
- Model deployment referenced by `AZURE_OPENAI_DEPLOYMENT_NAME` or `AZURE_FOUNDRY_MODEL_DEPLOYMENT`
- ToolSet containing: `get_expiring_contracts`, `get_vendor_spend_summary`, `search_cloud_blob_contracts`
- Instructions = full `SYSTEM_PROMPT` from `copilot_agent/agent.py`

What you must prepare first (on Azure)

- Foundry account / resource
- Foundry **project** (new Foundry project — not a classic hub-only workspace)
- At least one **chat model deployment** in that project
- Entra ID **RBAC** for the identity that runs the script
- Project **endpoint** URL copied into root `.env`
- (For tools at runtime) Fabric SQL + Azure AI Search reachable by the execution identity

PART A — Azure AI Foundry side (complete these first)

Work in this order. Each step ends with a **Done when** check so you know you can move on.

A1. Sign in and choose the right portal experience

1. Open <https://ai.azure.com> (Microsoft Foundry portal).
2. Sign in with an Azure account that can create resources in your target subscription.
3. If you see a **New Foundry** toggle, turn it **ON**. These steps refer to the current Foundry experience (not Foundry classic / hub-only).
4. Confirm the correct **Directory** and **Subscription** in the portal account menu.

Done when: You are on the Foundry home page for the intended subscription.

A2. Create the Foundry resource and project

You need a **Foundry project** under a Foundry / AI Services account. Prefer a **foundry-based project** (project endpoint visible on the overview). Classic hub-only projects often show only a connection string and may not work with this SDK path.

Option 1 — Portal (recommended)

1. In Foundry, choose **Create a new project** (or **Create** → Foundry / project).
2. Enter a **project name** (example: val-copilot-proj).
3. Open **Advanced** options if shown and set:
 - **Subscription**
 - **Resource group** (create new or reuse)
 - **Region** that supports **Agents** (examples commonly used: East US, West Europe, Southeast Asia — confirm Agents availability for your region)
 - Foundry / AI Services **resource name** if prompted
4. Create and wait until provisioning finishes (often several minutes).
5. Open the new project when ready.

Option 2 — Azure Portal

1. Azure Portal → **Create a resource** → search for **Azure AI Foundry / AI Services** (Foundry).
2. Create the account with a region that supports Agents.
3. From the account, create / open a **Project**.
4. Open the project in ai.azure.com.

Done when: Project Home / Overview loads and shows a **Project endpoint** (URL), not only a legacy hub connection string.

A3. Deploy a chat model in the project

The managed agent needs a **deployment name** that already exists in the project.

1. In the Foundry project, open **Discover** → **Models** (or **Deployments / Models + endpoints**, depending on UI).
2. Select a chat model suitable for tool-calling agents (example: gpt-4o or another GPT-4-class model available to your subscription).
3. Click **Deploy**.
4. Prefer **Custom settings** if you need a specific deployment name; otherwise **Default settings** is fine.
5. Record the **Deployment name** exactly (case-sensitive). Example: gpt-4o.
6. Wait until the deployment status is **Succeeded** / Ready.
7. Optional smoke test: open the model playground and send “hello” to confirm inference works.

Done when: Deployments list shows your model as ready, and you have copied the exact **deployment name**.

A4. Confirm Agents capability is available

1. In the project left navigation, open **Build** → **Agents** (or **Agents**).
2. Confirm the Agents blade loads (empty list is OK before the script runs).
3. If Agents is missing or blocked:
 - Switch to a **supported region**
 - Ensure you are on **New Foundry** (not classic hub-only)
 - Check subscription policy / quota with your Azure admin

Done when: You can open the Agents page for this project without an error.

A5. Assign Entra ID roles (RBAC)

The deploy script uses **DefaultAzureCredential** (your user, service principal, or managed identity). That identity must be allowed to create agents on the project.

1. Azure Portal → open the **Foundry account** or **project** resource.
 2. Open **Access control (IAM)** → **Add role assignment**.
 3. Assign at least one of:
 - **Azure AI Developer** (typical for agent create / run), or
 - **Azure AI Project Manager** / owner-equivalent if your org requires it
 4. Assign to the identity that will run `deploy_to_foundry.py` (your user UPN or the service principal).
 5. If teammates will use the agent, grant them the same (or Reader + run permissions per your org standard).
- Also plan data-plane access (needed when tools actually run):

Data system	Typical need
Microsoft Fabric SQL	Identity can connect via <code>ActiveDirectoryDefault</code> to the Gold warehouse
Azure AI Search	Search index reader (or API key in <code>.env</code> for key-based access)

Done when: The deploying identity has Foundry project rights, and Fabric/Search access is planned.

A6. Copy the Project endpoint (required for `.env`)

1. Open the Foundry project **Home / Overview / Welcome** page.
2. Find **Project endpoint** (copy button).
3. It should look like:


```
https://<ai-services-account-name>.services.ai.azure.com/api/projects/<project-name>
```
4. Paste this value into root `.env` as `AZURE_FOUNDRY_CONNECTION_STRING` (Part B).

Acceptable alternate form

Legacy semicolon connection string is also accepted by the script if it includes `endpoint=`:

```
endpoint=https://<account>.services.ai.azure.com/api/projects/<project>;subscriptionId=<guid>;resourceGroupName=<rg>
```

Warning — hub-only projects

If the portal only shows a classic **connection string** and **no** `.../api/projects/...` endpoint, you are likely on a hub-based project. Create a **new Foundry project** with a project endpoint before continuing.

Done when: You have saved the project endpoint URL and the model deployment name.

A7. Networking / firewall (if your tenant locks egress)

Complete this on Azure if corporate firewalls or private endpoints are enabled:

1. Foundry / Azure OpenAI networking blade → allow your runner's public IPs **or** private endpoint path.

2. Azure AI Search networking → allow the same runner / VNet.
 3. Fabric warehouse → allow the identity and network path used by ODBC.
- Done when:** A test call from the runner machine can reach Foundry (and later Fabric/Search).

A8. Azure Foundry readiness checklist (gate to Part B)

Do **not** run the Python deploy script until every box is checked:

- New Foundry portal experience enabled (not classic-only)
- Foundry account created in an Agents-supported region
- Foundry **project** created and openable in ai.azure.com
- **Project endpoint** URL copied ([https://...services.ai.azure.com/api/projects/...](https://...services.ai.azure.com/api/projects/))
- Chat **model deployed** and status = Ready
- Exact **deployment name** recorded (example: gpt-4o)
- **Agents** blade opens successfully
- Deploying identity has **Azure AI Developer** (or equivalent) on the project
- Fabric SQL + Azure AI Search access planned for tool runtime
- Networking allowlists completed (if required)

PART B — Code / script side (after Part A is complete)

B1. Clone the GitHub repository

```
git clone https://github.com/booshank/VALAgent.git
cd VALAgent
# use the branch that contains deploy_to_foundry.py if not yet on main
git checkout cursor/val-copilot-workspace-a9d1
```

B2. Configure root .env

Copy .env.example to .env at the repository root (do **not** create per-folder env files).

Map values collected in Part A:

Variable	Required	Source (Azure side)
AZURE_FOUNDRY_CONNECTION_STRING	Yes	Project endpoint from A6
AZURE_FOUNDRY_MODEL_DEPLOYMENT or AZURE_OPENAI_DEPLOYMENT_NAME	Yes	Exact deployment name from A3
AZURE_FOUNDRY_AGENT_NAME	No	Desired agent name (default val-copilot)
FABRIC_SQL_SERVER / FABRIC_SQL_DATABASE	Runtime	Fabric warehouse
AZURE_SEARCH_ENDPOINT / API key / index	Runtime	Azure AI Search

Example:

```
AZURE_FOUNDRY_CONNECTION_STRING=https://myacct.services.ai.azure.com/api/projects/val-copilot-proj
AZURE_FOUNDRY_MODEL_DEPLOYMENT=gpt-4o
AZURE_FOUNDRY_AGENT_NAME=val-copilot
AZURE_OPENAI_DEPLOYMENT_NAME=gpt-4o
```

B3. Authenticate (DefaultAzureCredential)

No Foundry API key is used by the deploy script for control-plane auth.

1. **Developer workstation:** `az login` then `az account set --subscription <id>`
2. **Service principal:** set `AZURE_CLIENT_ID`, `AZURE_TENANT_ID`, `AZURE_CLIENT_SECRET`
3. **CI / VM:** managed identity with the RBAC from **A5**

```
az login
az account show --query id -o tsv
```

B4. Install Python dependencies

```
cd copilot_agent
python3 -m venv .venv
source .venv/bin/activate
pip install -r requirements.txt
pip install -r ../mcp_server/requirements.txt
```

Pinned packages: `azure-ai-projects==1.0.0`, `azure-ai-agents>=1.0.0,<2.0.0`,
`azure-identity>=1.19.0`. Do **not** upgrade to `azure-ai-projects 2.x` for this
script — ToolSet was removed.

B5. Dry-run (local validation only)

```
cd copilot_agent
python deploy_to_foundry.py --dry-run -v
```

Expected:

- status: `dry_run`
 - tools: `get_expiring_contracts`, `get_vendor_spend_summary`, `search_cloud_blob_contracts`
 - non-zero `instructions_chars`
 - endpoint parsed from `AZURE_FOUNDRY_CONNECTION_STRING`
-

B6. Create the managed agent in Foundry

```
cd copilot_agent
python deploy_to_foundry.py -v
# optional:
python deploy_to_foundry.py --agent-name val-copilot-prod
```

What the script does

1. Load credentials via `DefaultAzureCredential`
2. Parse `AZURE_FOUNDRY_CONNECTION_STRING` → `AIPProjectClient`
3. Import MCP tools from `mcp_server/server.py`
4. Wrap tools with `FunctionTool` + `ToolSet`
5. Call `project_client.agents.create_agent(...)` with `model`, `name`, `instructions=SYSTEM_PROMPT`, `toolset`
6. Print JSON including `agent_id` on success

This is the step that creates the agent object inside your Foundry project. Part A only prepared the project/model; Part B registers VAL CoPilot.

B7. Verify on the Azure Foundry side (after script success)

Return to `ai.azure.com` → your project:

1. Open **Build** → **Agents**.
2. Confirm agent `val-copilot` (or your custom name) appears.
3. Open the agent and check:
 - **Instructions** contain `Comparative Analysis + Contract Lifecycle Procedures`

- **Tools** list the three function tools

- **Model** matches the deployment from **A3**

4. Use the **Agent playground** (if available) with:

- “Show expiring contracts”

- “Red-flag compliance audit for CON-0003”

5. If tool calls fail in playground, verify Fabric/Search credentials and that the host executing function tools can reach those services.

Runtime & architecture notes

The Foundry-managed agent stores **instructions + tool definitions**. Function tool **implementations** are the Python MCP functions from this repo and run in the process that handles tool calls. Ensure that host can reach Fabric SQL (ODBC Driver 18 + ActiveDirectoryDefault) and Azure AI Search.

Cognitive procedures encoded in SYSTEM_PROMPT

- Comparative Analysis & Decision Framework → ## Recommendation
- Red-Flag Compliance Audits → ## Red-Flag Compliance Audit
- Dynamic Counter-Clause Drafting → ## Dynamic Counter-Clause Drafting
- Financial Exposure Projections → ## Financial Exposure Projection
- Proactive Renewal Strategy Sheets → ## Proactive Renewal Strategy Sheet

Troubleshooting

Symptom	Likely cause / fix
No Project endpoint in portal	Hub-only / classic project — create a new Foundry project (A2)
Agents blade missing	Unsupported region or classic experience — change region / enable New Foundry (A4)
Missing AZURE_FOUNDRY_CONNECTION_STRING	Paste endpoint from A6 into root .env
create_agent auth / 403	RBAC missing — assign Azure AI Developer (A5); re-login
Wrong model / deployment not found	Deployment name mismatch — copy exact name from A3
Cannot import ToolSet / FunctionTool	Reinstall pinned azure-ai-projects==1.0.0 and azure-ai-agents<2
Tool returns empty / SQL errors	Fabric/Search env or network (A7); test MCP locally
Firewall / VNet blocks	Allowlist runner IPs on Foundry + Search (A7)

Full deployment checklist

Azure Foundry side (Part A)

- A1 Portal sign-in + New Foundry ON
- A2 Foundry project created
- A3 Model deployed (name recorded)
- A4 Agents blade opens
- A5 RBAC assigned
- A6 Project endpoint copied
- A7 Networking done (if needed)
- A8 Readiness gate passed

Code / script side (Part B)

- Repo cloned
- Root .env populated from Part A values
- az login / managed identity ready
- Dependencies installed
- --dry-run succeeds
- deploy_to_foundry.py returns status=created + agent_id
- Portal Agents view shows agent, tools, and SYSTEM_PROMPT procedures
- Playground / smoke-test prompts succeed

Keep secrets out of Git; commit only .env.example. Script path: copilot_agent/deploy_to_foundry.py.

VALAgent · Keep secrets out of Git · Script: copilot_agent/deploy_to_foundry.py